




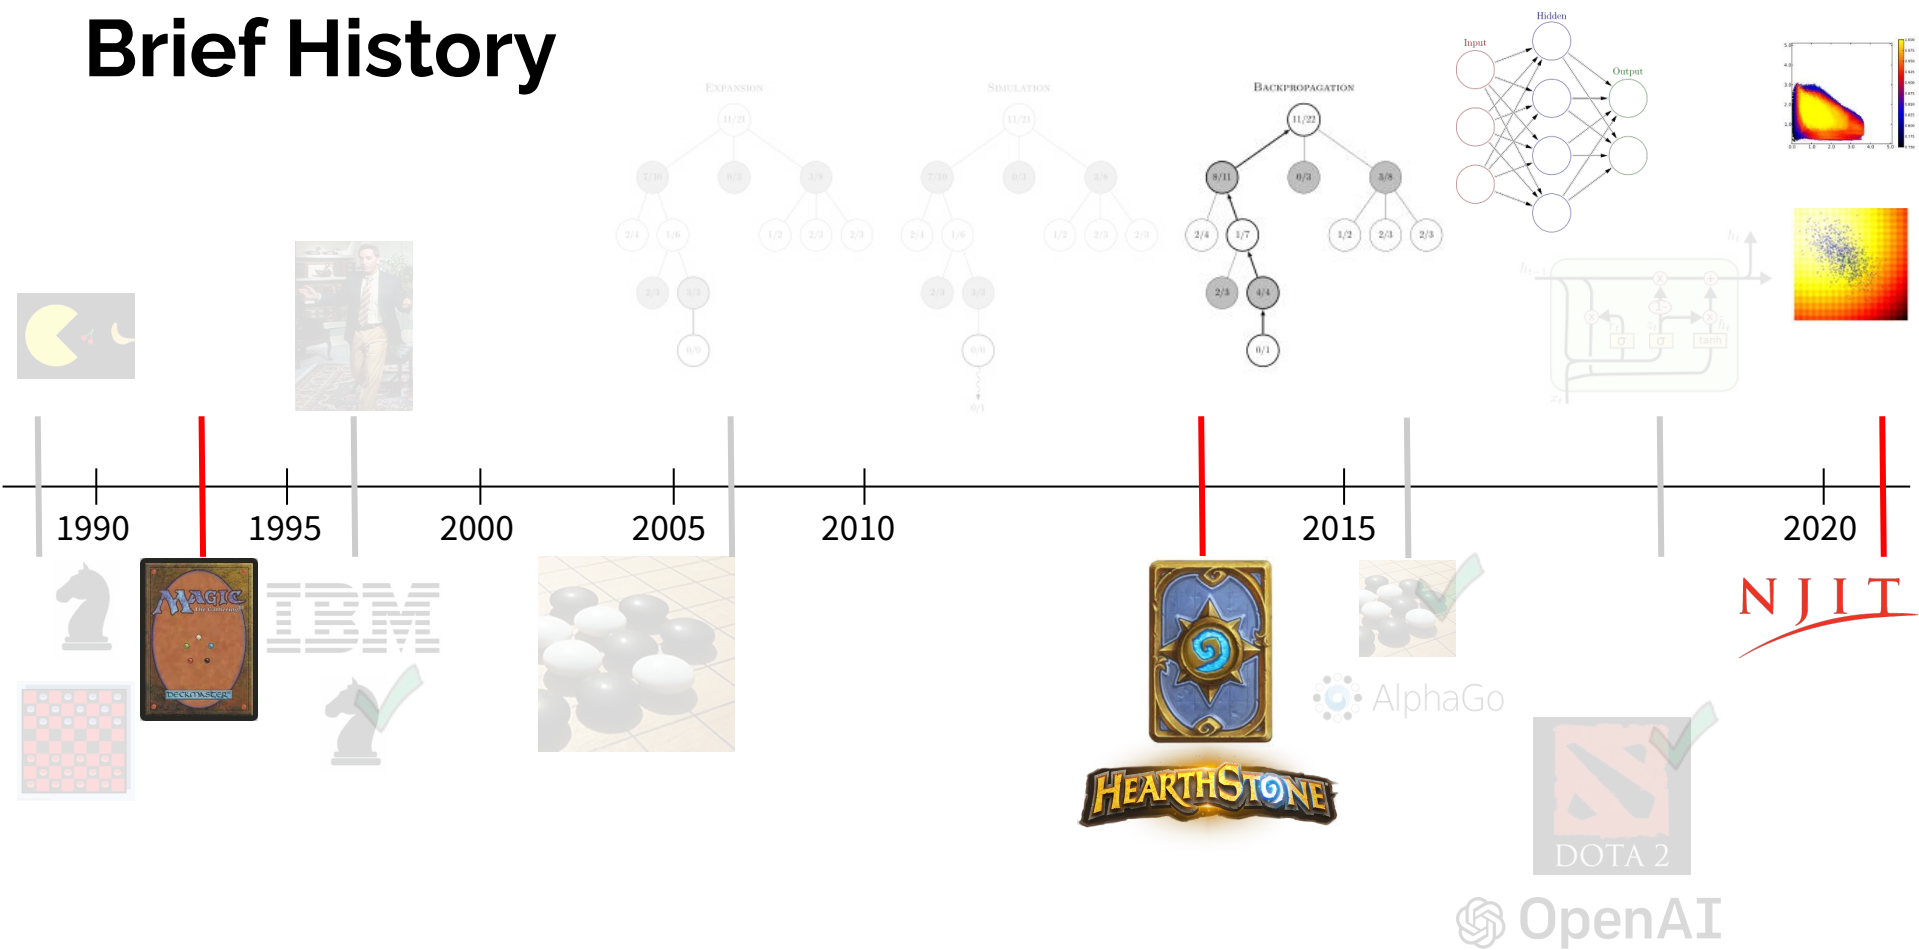
Analysis of Gameplay Strategies in Hearthstone: A Data Science Approach

Connor W. Watson | Dr. Amy K. Hoover | Dr. Usman W. Roshan | Dr. Senjuti Basu Roy

Brief History



Brief History



Why Games?

- Hard to “predict”:
 - Chess has 6 pieces, Hearthstone over 3000
 - Multiple strategies (5) + sub-strategies (100?)
 - Observability
 - Chess (deterministic) vs Hearthstone (stochastic) [1]
- Hearthstone has orders of magnitude larger solution space than Chess/Go
 - Game tree + level of uncertainty grow (MCTS) ❌
 - Brute force is challenging - shown for Backgammon, Checkers, Chess ❌
 - Need a better way to make decisions

Abstract


High Level Goals:

1. Compare/contrast & visualize gameplay strategies
2. Use new algorithm to make better players
 - a. Can agents make better decisions
3. Predict gameplay strategies

What is ~~Hearthstone~~ Chess?

Remove the pieces...

Leave the King on board





What is ~~Hearthstone~~ Chess?

Pick 30

Bring those to
games

King has abilities

Where are the
pieces...




What is ~~Hearthstone~~...Chess?

Hide pieces

Every turn you put some

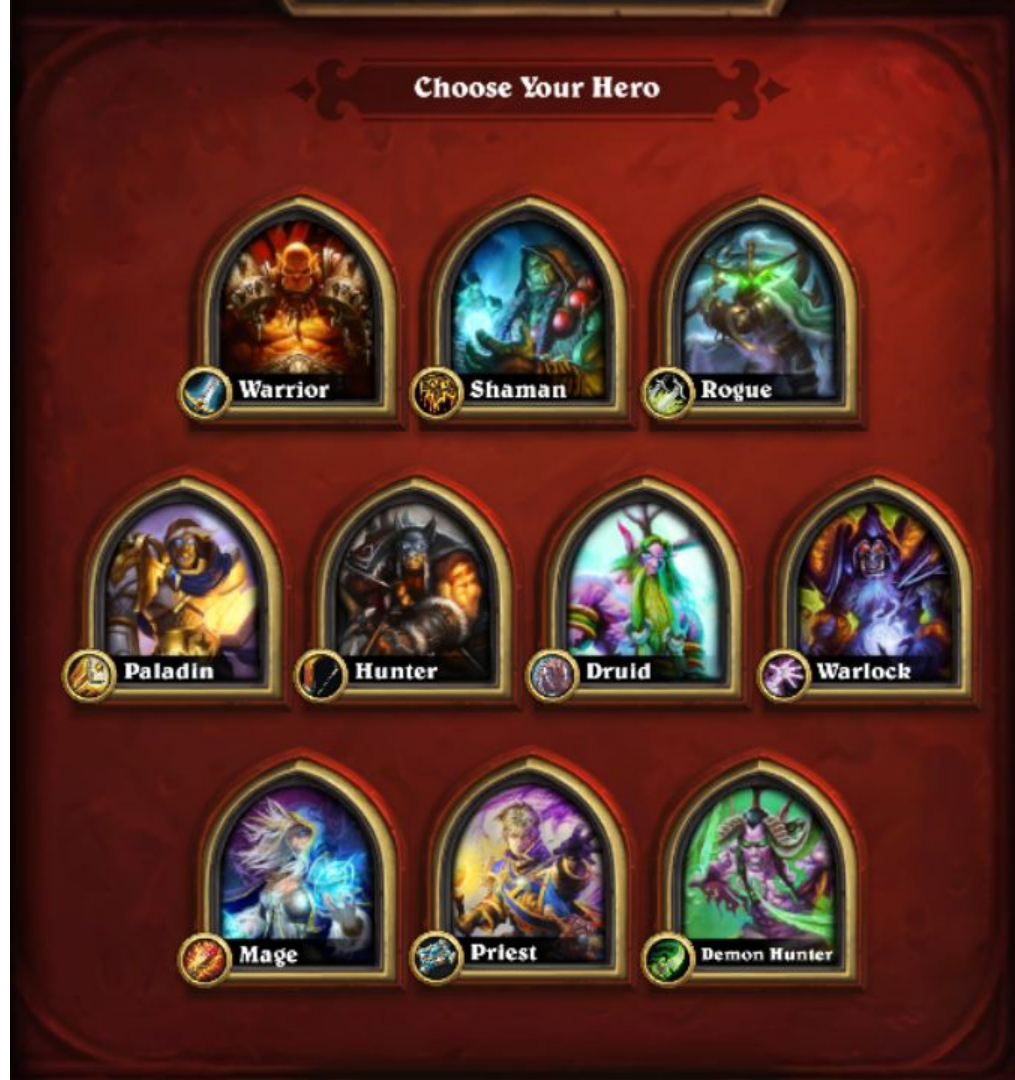
What's the goal?

Goal: hit king 30 times



Hero Classes

- 10 to choose from
- Each has specific cards
- “Basic” cards - shared with all




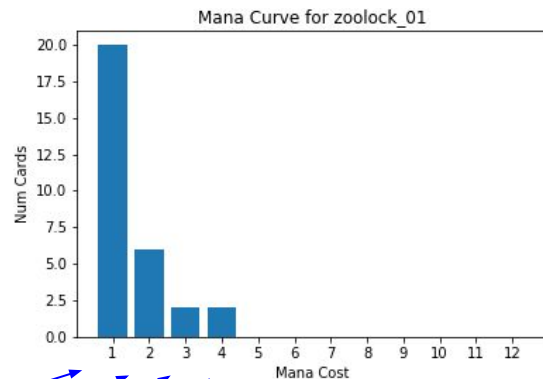
Types of Cards - Minions/Spells

- Mana cost (top left)
 - Name (center) /Effect (lower middle)
 - Attack power (bottom left)
 - Health points (bottom right)
 - Played onto the board
- Mana cost (top left)
 - Name (center) /Effect (lower middle)
 - Played from hand (single use)
 - Not on the board



Deck Building Strategies

- Mana Curve - amount of cards in each bin of mana cost
 - Resource management tied to strategy
- Aggro deck (top)
 - Many low cost cards
- Control deck (bottom)
 - Scattered across all bins




SabberStone: a HS Simulator

- Hearthstone is not open source
- SabberStone is a community repository
 - GitHub: <https://github.com/HearthSim/SabberStone>
- Comes packaged with AI for research / testing
 - Turn local game tree
 - Heuristic scoring functions

Heuristic Scoring Functions

- AggroScore
 - Rewards attacking opponent
 - End game faster
- ControlScore
 - Rewards board control
 - End game slower
- Can an algorithm evolve better scoring functions?



Heuristic Scoring Functions

```
//ControlScore
public override int Rate()
{
    if (OpHeroHp < 1)
        return int.MaxValue;
    if (HeroHp < 1)
        return int.MinValue;
    int result = 0;
    if (OpBoardZone.Count == 0 && BoardZone.Count > 0)
        result += 1000;

    result += (BoardZone.Count - OpBoardZone.Count) * 50;
    result += (MinionTotHealthTaunt - OpMinionTotHealthTaunt) * 25;
    result += (HeroHp - OpHeroHp) * 10;

    result += MinionTotAtk;
    return result;
}
```

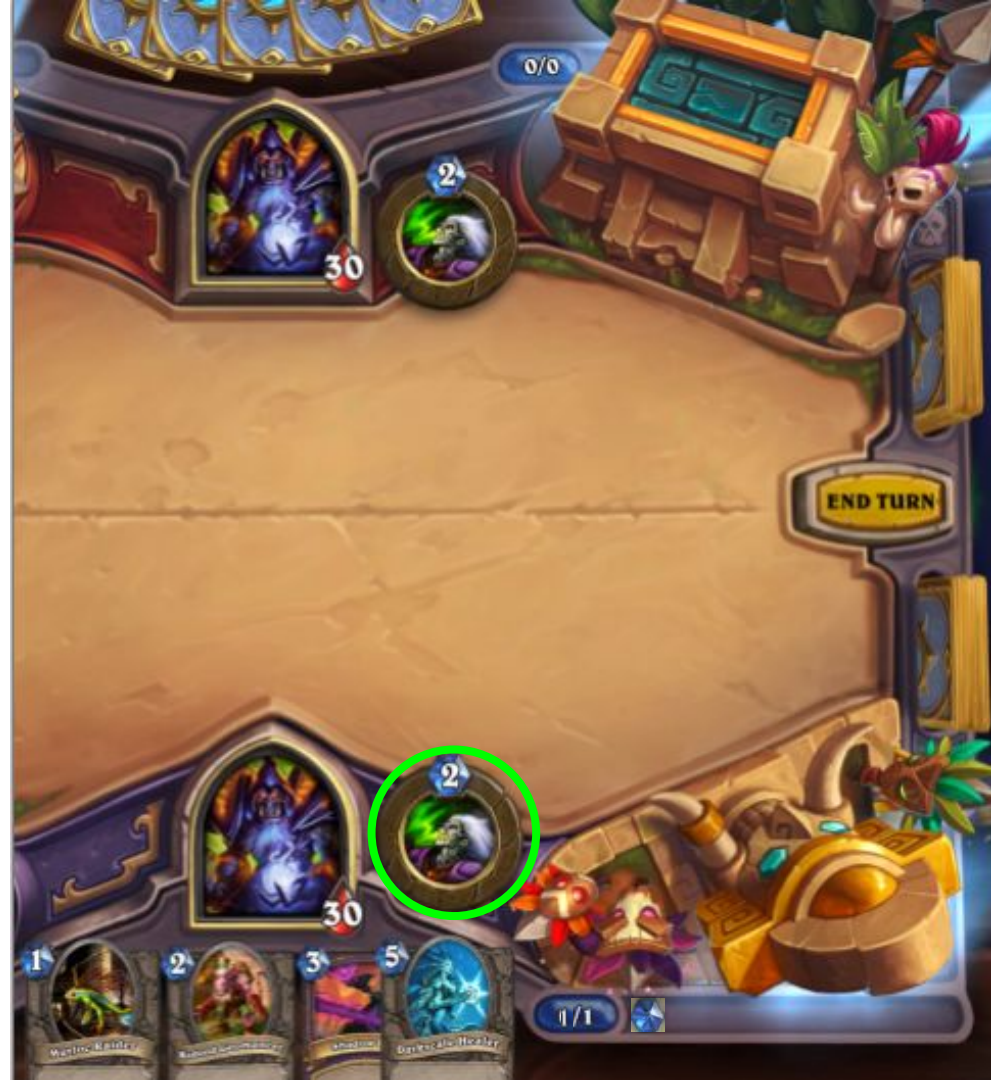
```
//AggroScore
public override int Rate()
{
    if (OpHeroHp < 1)
        return int.MaxValue;
    if (HeroHp < 1)
        return int.MinValue;
    int result = 0;
    if (OpBoardZone.Count == 0 && BoardZone.Count > 0)
        result += 1000;

    if (OpMinionTotHealthTaunt > 0)
        result += OpMinionTotHealthTaunt * -1000;
    result += (HeroHp - OpHeroHp) * 1000;

    result += MinionTotAtk;
    return result;
}
```

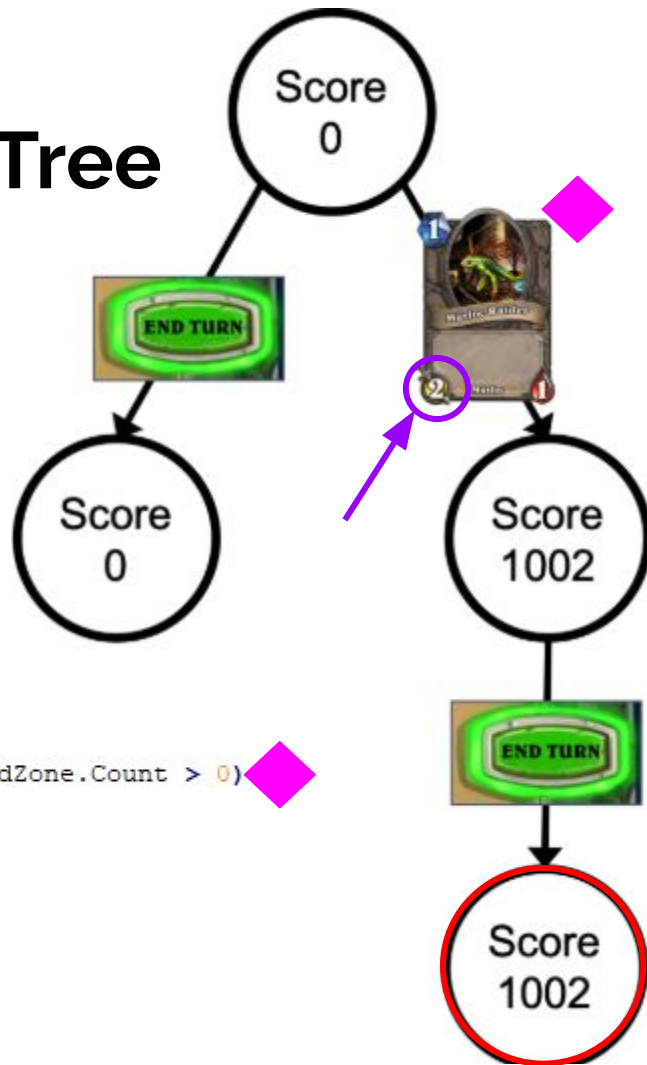

Game States

- Starting turn 1 (bottom)
- Both players start with 0
- Gain 1 mana at start of turn
- 1 mana to spend
 - 1 cost cards only



Game States/Tree

- Most positive score
- Best decision

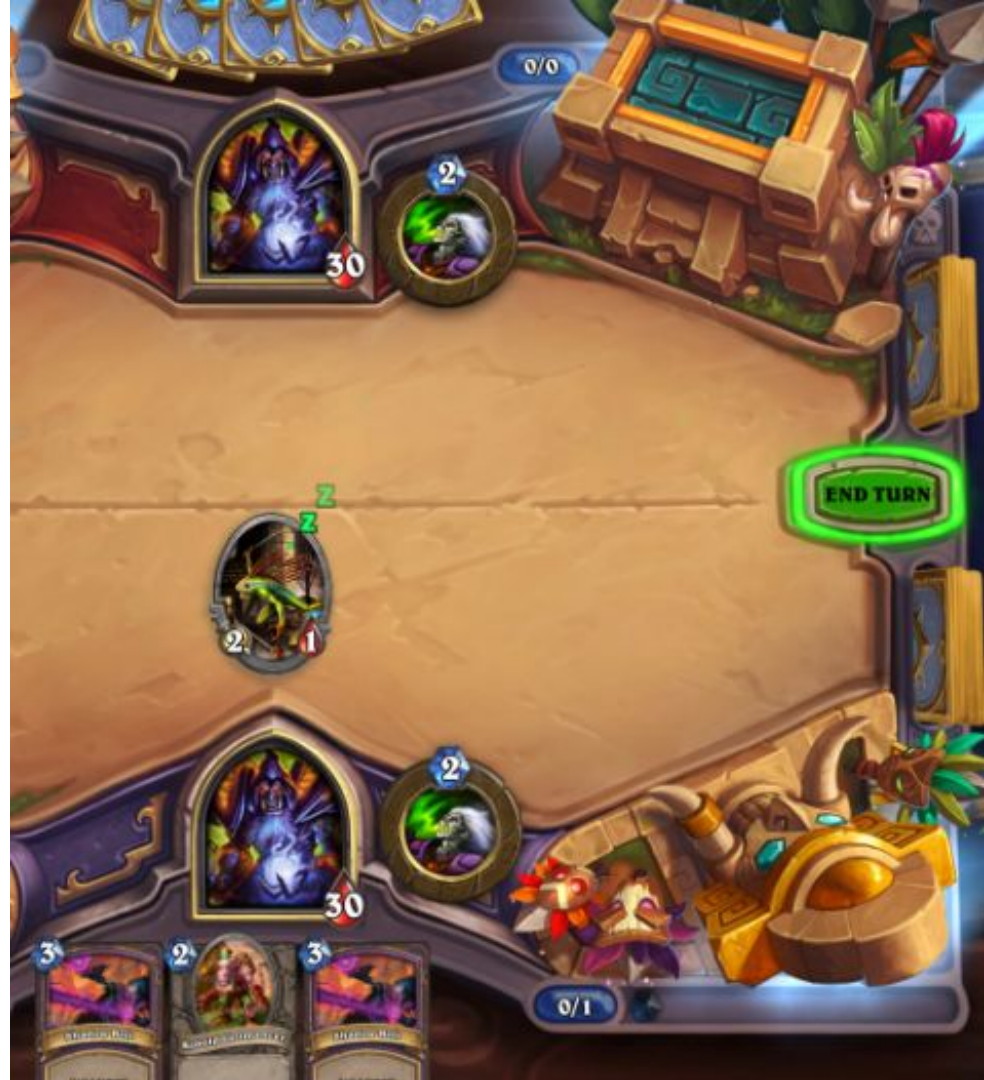


```
int result = 0;  
if (OpBoardZone.Count == 0 && BoardZone.Count > 0) {  
    result += 1000;  
    ...  
    result += MinionTotAtk;   
    return result;  
}
```



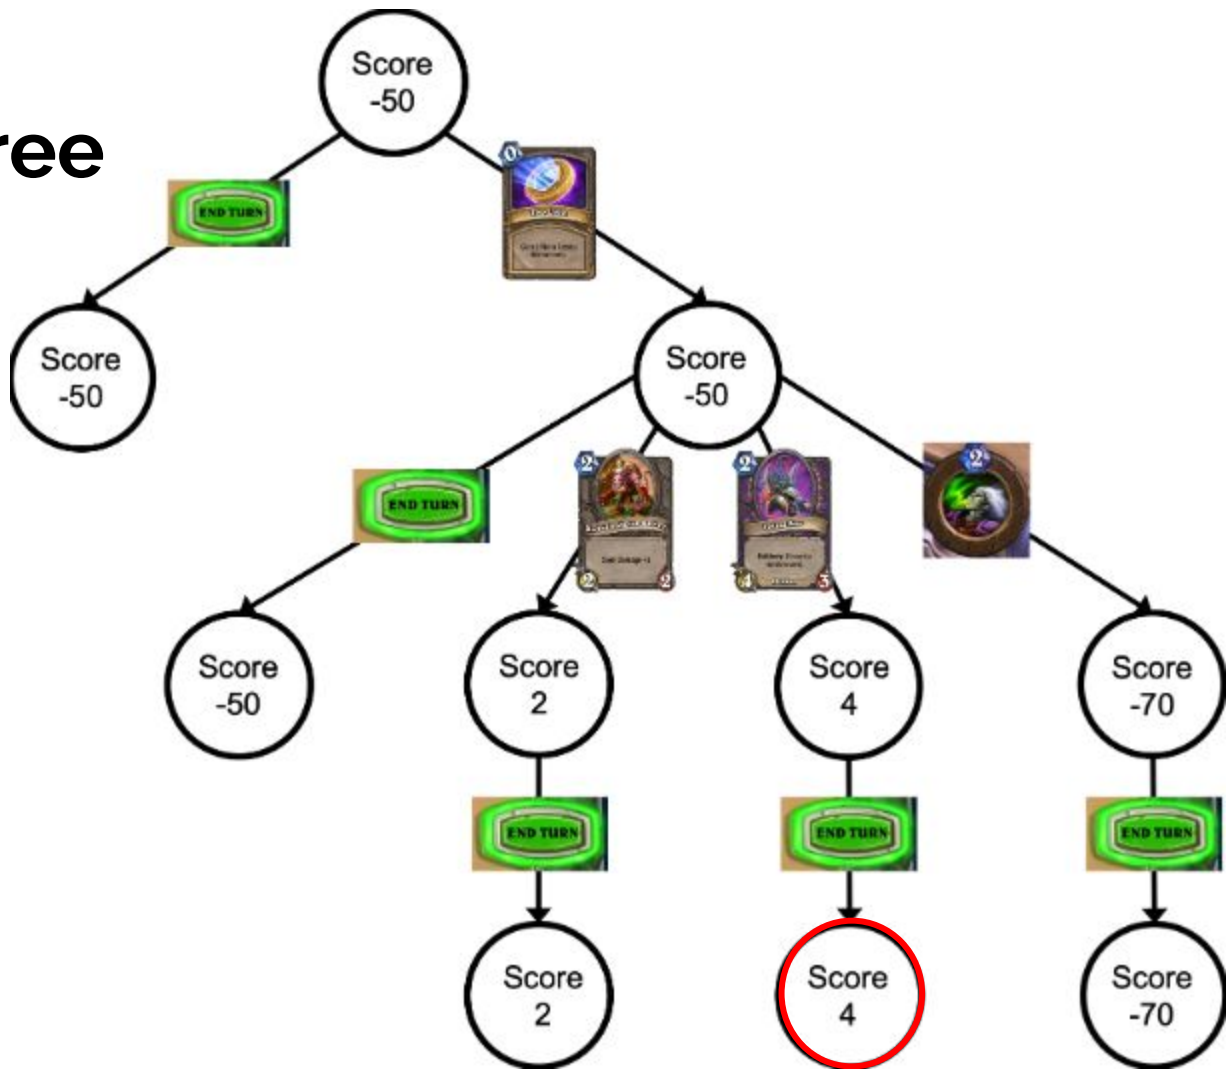
Game States

- Choice: play minion card
- Maintain board control




Game States/Tree

- Tree ends at current turn
- No look-ahead
- Start at disadvantage



CMA-ES

- Calculate fitness for generation (solutions)
- Isolate top N %
- Calculate covariance matrix of next generation
 - Uses covariance matrix only
- Sample new solutions from it



Source [7]


MAP-Elites

- Average mana cost = 2.5
- Variance = 1.25
- Behavior Vector: (2.5, 1.25)
- Fitness (win rate) = 100/200
 - NumGamesWon / NumGamesPlayed



D1: F=100

Map of Elites



Buffer


D1: F=100

MAP-Elites

- Average mana cost = 2.5
- Variance = 1.25
- Behavior Vector: (2.5, 1.25)
- Fitness (win rate) = 100/200
 - NumGamesWon / NumGamesPlayed



D1: F=100



Buffer

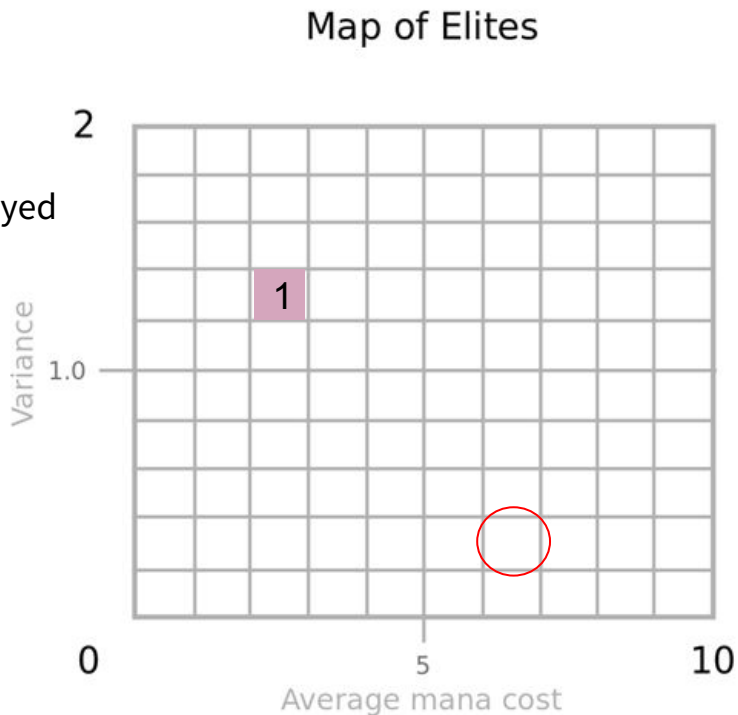
D1: F=100

MAP-Elites

- Average mana cost = 6.5
- Variance = 0.25
- Behavior Vector: (6.5, 0.25)
- Fitness (win rate) = 80/200
 - NumGamesWon / NumGamesPlayed



D2: F=80



Buffer

D1: F=100

D2: F=80


MAP-Elites

- Average mana cost = 6.5
- Variance = 0.25
- Behavior Vector: (6.5, 0.25)
- Fitness (win rate) = 80/200
 - NumGamesWon / NumGamesPlayed



D2: F=80

Map of Elites



Buffer

D1: F=100

D2: F=80

MAP-Elites


- Average mana cost = 2.5
- Variance = 1.25
- Behavior Vector: (2.5, 1.25)
- Fitness (win rate) = 180/200
 - NumGamesWon / NumGamesPlayed



What happens?



D3: F=180



Buffer

D1: F=100

D2: F=80


D3: F=180

MAP-Elites

- Average mana cost = 2.5
- Variance = 1.25
- Behavior Vector: (2.5, 1.25)
- Fitness (win rate) = 150/200
 - NumGamesWon / NumGamesPlayed



D3: F=180



Buffer




D1: F=100

D2: F=80

D3: F=180

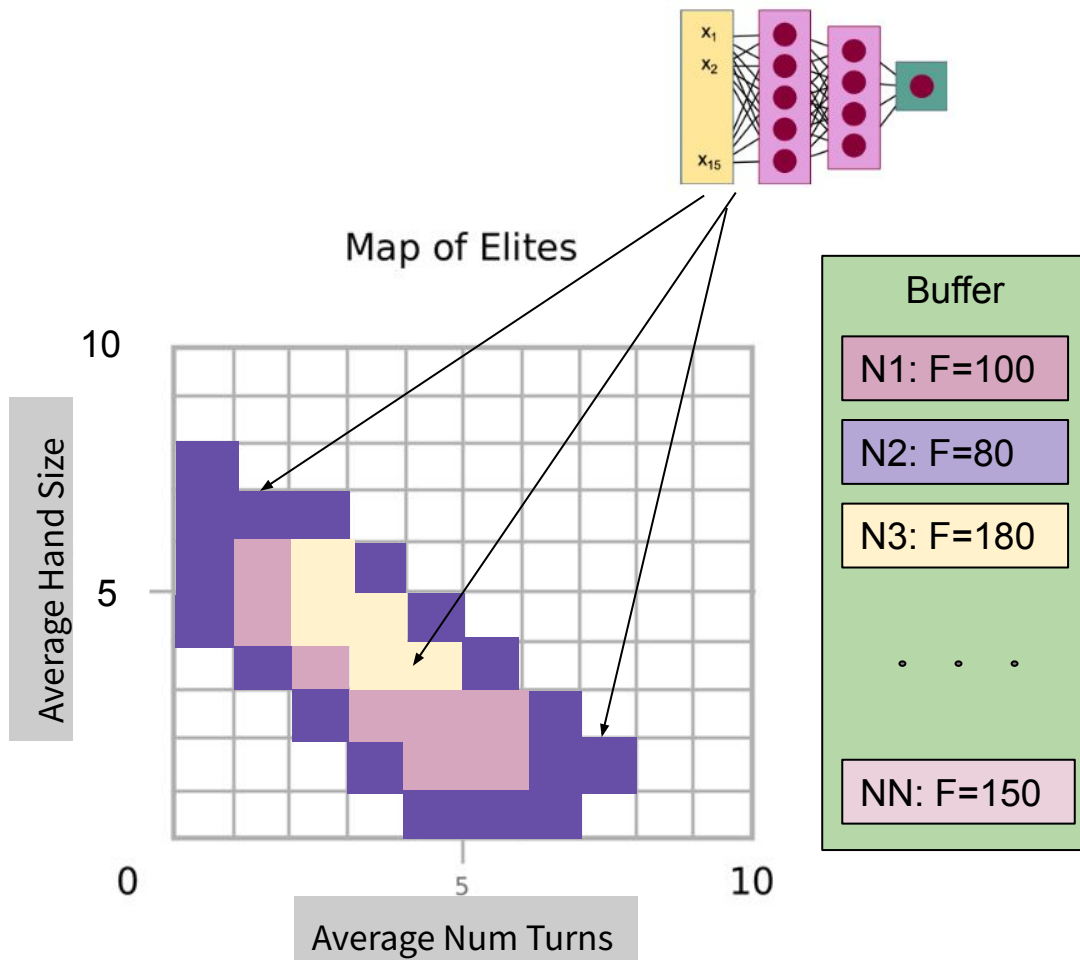

MAP-Elites

- Keep the “Elite” decks (solutions)
- These have the highest win rate in each grid cell



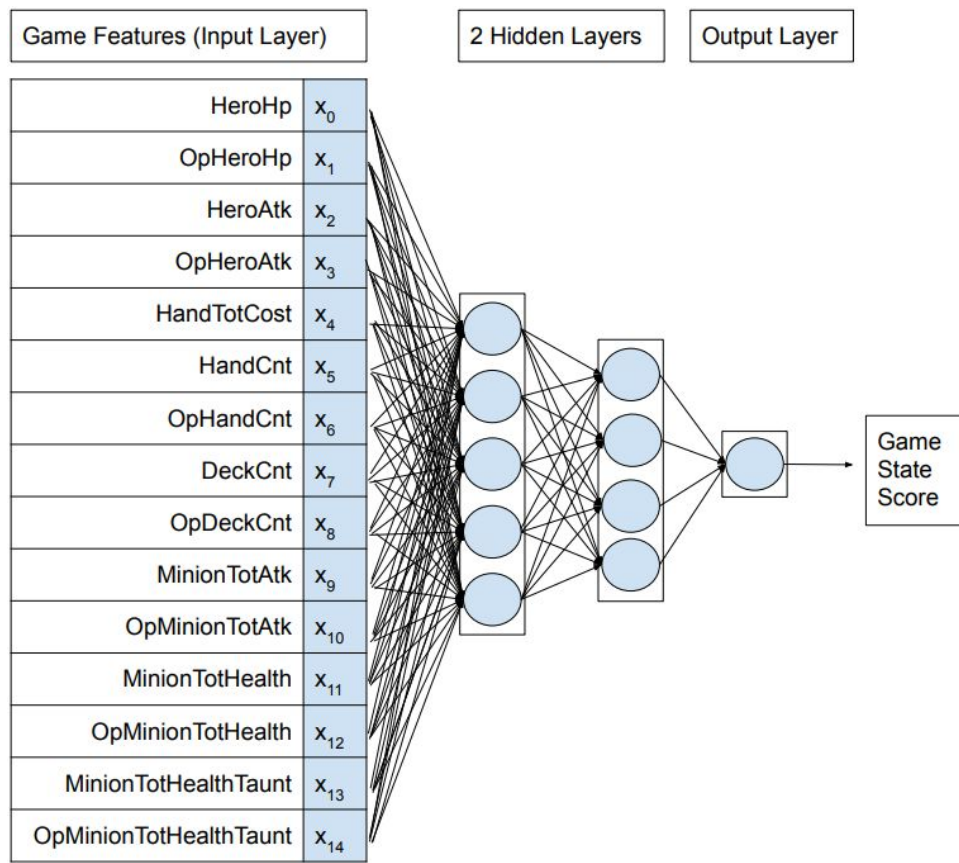
CMA-ME

- Use ANNs instead of decks
- Evolve with CMA-ES
- MAP of Elite ANNs
- Play 200 games with 5000 ANNs



Artificial Neural Network

- Fully connected
 - All nodes connect from one layer to next
- Feed forward
 - No backward connections
- Inspired by [8, 9]
- Input:
 - Visible game pieces
- Output:
 - Game state score



Research Questions

1. Is control really better than aggro?
 - a. Or is ControlScore a better heuristic
2. Can ANNs perform better?
3. Can gameplay strategies be predicted?
4. Can gameplay strategies be visualized?



The Data

25 total columns...

- Collection of turns (all games)
- Num turns per game, which player won, etc..
- How to reduce this?

| TURN_NO | P1_HEALTH | P2_HEALTH | CURRENT_PLAYER | AMOUNTHEALEDTHISTURN |
|---------|-----------|-----------|------------------------|----------------------|
| 1 | 30 | 30 | P1 FitzVonGerald | 0 |
| 2 | 30 | 30 | P2 RehHausZuckFuchs | 0 |
| 3 | 30 | 29 | P1 FitzVonGerald | 0 |
| 4 | 27 | 29 | P2 RehHausZuckFuchs | 0 |
| 5 | 27 | 24 | P1 FitzVonGerald | 0 |
| 6 | 21 | 24 | P2 RehHausZuckFuchs | 0 |
| 7 | 21 | 16 | P1 FitzVonGerald | 0 |
| 8 | 8 | 16 | P2 RehHausZuckFuchs | 0 |
| 1 | 30 | 30 | P1 FitzVonGerald | 0 |

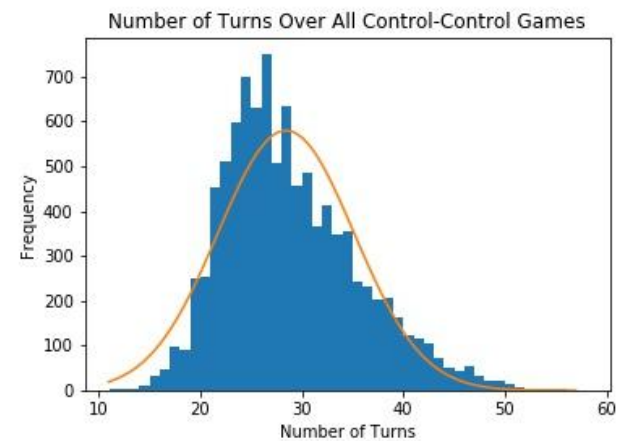
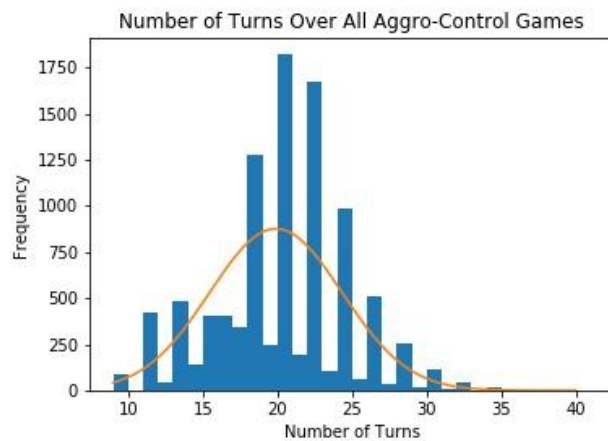
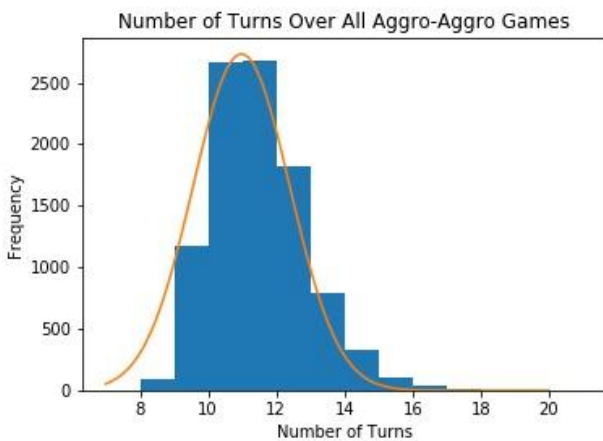
The Data

13 total columns...

- Lots of feature engineering
- Player statistics per turn/game

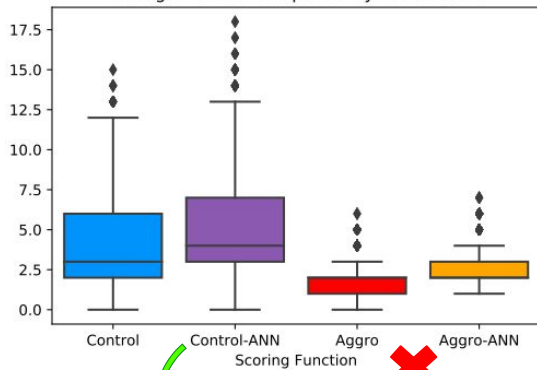
| PlayerStrategy | AvgHealedPerTurn | AvgHeroAttacksPerTurn | AvgCardsDrawnPerTurn | AvgCardsPlayedPerTurn |
|----------------|------------------|-----------------------|----------------------|-----------------------|
| 0.0 | 1.17 | 0.0 | 1.50 | 1.50 |
| 0.0 | 0.00 | 0.0 | 1.14 | 0.86 |
| 0.0 | 0.00 | 0.0 | 1.00 | 1.17 |
| 0.0 | 0.00 | 0.0 | 1.50 | 1.00 |
| 0.0 | 0.00 | 0.0 | 1.00 | 1.40 |
| 1.0 | 0.83 | 0.0 | 1.17 | 1.00 |
| 1.0 | 0.36 | 0.0 | 1.91 | 1.36 |
| 1.0 | 0.33 | 0.0 | 1.17 | 1.00 |
| 1.0 | 0.36 | 0.0 | 2.09 | 1.18 |
| 1.0 | 0.14 | 0.0 | 1.21 | 1.21 |

The Data




The Data

Average Number of Spells Played Per Game




The Data

Average Number of Spells Played Per Game




Average Number of Friendly Minion Attacks Per Turn




The Data


Average Number of Spells Played Per Game



Average Number of Friendly Minion Attacks Per Turn



Number of Hero Power Activations Per Game




Experiment 1

- Test aggro players vs control players
- Test decks using opposite scoring function
- Which performs better?
 - Is control > aggro generally?
 - Or is ControlScore > AggroScore?
- Hypotheses:
 - $\bar{X}_{\text{control}} > \bar{X}_{\text{aggro}}$
 - $\bar{X}_{\text{ControlScore}} > \bar{X}_{\text{AggroScore}}$


Results

- ControlScore > AggroScore
 - bottom
- 50/50 in mirror matches
 - (right)


Win Rates for AggroScore&Deck
vs ControlScore&Deck



Win Rates for ControlScore&Deck
vs ControlScore&Deck




Win Rates for AggroScore&Deck
vs AggroScore&Deck




Results

- ControlScore > AggroScore
 - ControlScore won in all three charts


Win Rates for ControlScore&Deck
vs AggroScoreControlDeck



Win Rates for ControlScoreAggroDeck
vs AggroScoreControlDeck



Win Rates for AggroScore&Deck
vs ControlScoreAggroDeck



Experiment 2

- Evolve networks with different behavior features
 - Number of ANNs to evolve
 - Change num turns / hand size
- Hypothesis:
 - $\bar{x}_{ANN} > \bar{x}_{Sabber}$



Table 4.3 CMA-ME Behavior Configurations^a

| <i>Network Name</i> | <i>Num Games per ANN</i> | <i>Num ANNs To Evaluate</i> | <i>HandSize [Min,Max]</i> | <i>Num Turns [Min,Max]</i> | <i>PlayerScore, OpponentScore</i> |
|---------------------|--------------------------|-----------------------------|---------------------------|----------------------------|-----------------------------------|
| Warlock Net_CC_sm | 100 | 5000 | [1,7] | [5,15] | Control,Control |
| CvsNNC_2.0 | 100 | 5000 | [1,7] | [25,35] | Control,Control |
| CvsNNC_Large | 200 | 50000 | [1,9] | [5,45] | Control,Control |
| Warlock Net_AA_sm | 100 | 5000 | [1,7] | [5,15] | Aggro,Aggro |
| Warlock Net_AA_lg | 200 | 50000 | [1,9] | [5,45] | Aggro,Aggro |







Results

- Aggro networks (left) look much different than control networks (right)
- Main difference:
 - num turns (x axis)
- Where are the ANNs focused?




Results

- Aggro networks (left) look much different than control networks (right)
- Changed behaviors:
 - NumTurns [5,15] -> [5,45]
 - NumGames 100->200
- How do these change ANNs?



Results


- Recall AggroScore v AggroScore had roughly 50% win rate
- ANN aggro players performed better than AggroScore




Results

- Top: same configs as aggro evolution
- Bottom: expanded num turns / hand size
- Change behaviors ->
control ANNs perform better
 - Particularly for hand size / num turns
 - Evolves a stronger control network

Win Rates for ControlScore
vs EvoControlScore



Win Rates for ControlScore
vs EvoControlScore



Experiment 3

- Can a model predict / **generalize** gameplay strategy?
- Compare across **five** models [6]
- Train using aggro and control players
 - AggroScore/ControlScore
- Test using players with ANN heuristics

Results

- High accuracy on train/validation data
- Lower accuracy on test data
- Precision > Recall

Validation Data - Aggro/ControlScores

| <i>Model</i> | <i>Accuracy</i> | <i>Precision</i> | <i>Recall</i> |
|---------------------|-----------------|------------------|---------------|
| Logistic Regression | 0.9986 | 0.9981 | 0.9992 |
| Random Forest | 0.9990 | 0.9988 | 0.9992 |
| SVM | 0.9987 | 0.9983 | 0.9992 |
| Decision Tree | 0.9981 | 0.9978 | 0.9983 |
| SGD Classifier | 0.9988 | 0.9985 | 0.9992 |

Test Data - ANN Heuristics



| <i>Model</i> | <i>Accuracy</i> | <i>Precision</i> | <i>Recall</i> |
|---------------------|-----------------|------------------|---------------|
| Logistic Regression | 0.6106 | 0.7390 | 0.3418 |
| Random Forest | 0.6623 | 0.8201 | 0.4157 |
| SVM | 0.6120 | 0.7408 | 0.3446 |
| Decision Tree | 0.6817 | 0.7854 | 0.4999 |
| SGD Classifier | 0.6245 | 0.7708 | 0.3543 |

Results

- Control class mostly **correct**
- Aggro class mostly **incorrect**
 - Strategies tend to “bleed”
- Aggro-ANN features may fall into the ControlScore distribution


SVM Classifier Confusion Matrix

| | Actual | | | |
|-----------|---------|-------|---------|-------|
| | | Aggro | Control | |
| Predicted | Aggro | 3532 | 1201 | 4733 |
| | Control | 6526 | 8757 | 15283 |
| | | 9958 | 9958 | 19916 |





Experiment 4


- How to visualize players in Cartesian space?
- PCA for space reduction
 - Projected on 2 axes
 - Aggro/Control Score only



Results




PCA Biplot of Averaged Game Stats per Turn



Results


- Support Vector Classifier
- 98% accurate on validation data




Results

- 75% accurate on test data (ANNs) vs 61% before

Number of Hero Power Activations Per Game




Average Number of Friendly Minion Attacks Per Turn



PC2: [17.58]

PCA Score Plot of Averaged Game Stats per Turn



Conclusions

- What **is** validated
 - ControlScore better than AggroScore
 - CMA-ME produces better heuristics
 - Other applications?
- What **needs to be** improved
 - Supervised learning models
- Future work
 - More heroes / decks
 - Human players
 - Turn-by-turn prediction



Thank You

- Dr. Hoover
- Dr. Roshan
- Dr. Basu Roy
- The Public

Questions



Bibliography

- [1] Stiegler, Andreas, et al. "Symbolic reasoning for hearthstone." IEEE Transactions on Games 10.2 (2017): 113-127.
- [2] M. Campbell, A. J. Hoane Jr, and F.-h. Hsu, "Deep blue," Artificial intelligence, vol. 134, no. 1-2, pp. 57–83, 2002.
- [3] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot et al., "Mastering the game of go with deep neural networks and tree search," nature, vol. 529, no. 7587, p. 484, 2016.
- [4] M. C. Fontaine, J. Togelius, S. Nikolaidis, and A. K. Hoover, "Covariance matrix adaptation for the rapid illumination of behavior space," arXiv preprint arXiv:1912.02400, 2019.
- [5] A. Bhatt, S. Lee, F. de Mesentier Silva, C. W. Watson, J. Togelius, and A. K. Hoover, "Exploring the hearthstone deck space," in Proceedings of the 13th International Conference on the Foundations of Digital Games. ACM, 2018, p. 18.
- [6] Fernández-Delgado, Manuel, et al. "Do we need hundreds of classifiers to solve real world classification problems?." The journal of machine learning research 15.1 (2014): 3133-3181.
- [7] D. Ha, "A visual guide to evolution strategies," blog.otoro.net, 2017. [Online].Available: <https://blog.otoro.net/2017/10/29/visual-evolution-strategies/>
- [8] Cuccu, Giuseppe, Julian Togelius, and Philippe Cudré-Mauroux. "Playing atari with six neurons." Proceedings of the 18th international conference on autonomous agents and multiagent systems. International Foundation for Autonomous Agents and Multiagent Systems, 2019.
- [9] Bellemare, Marc G., et al. "The arcade learning environment: An evaluation platform for general agents." Journal of Artificial Intelligence Research 47 (2013): 253-279.

Appendix (Unused)

Appendix should contain supporting charts / slides with more details. This section was unused. Instead, some draft slides were pushed to the back.

Hyperparameter Search

- SVM / LogReg -
 - l1/l2 regularization,
 - C (0.1, 1, 10, 100) - regularization parameter
- RF number of estimators (50, 100, 200)
- DT
 - (gini, entropy)
 - splits (best, random)
 - max depth (4-10)
- SGD loss (hinge, log, perceptron, modified huber),
 - l1/l2 regularization
 - learning rate (1, 0.1, 0.01, 0.001)

Hyperparameter Search Results

The results of the hyperparameter search are below:

- Logistic Regression: l2 regularizer, C coefficient of 10
- Random Forest: 100 estimators
- Support Vector Classifier: l2 regularizer, C coefficient of 1
- Decision Tree Classifier: best splitters, max depth of 10, entropy criterion
- SGD Classifier: modified huber loss, l2 regularizer, alpha 0.001